



Hochschule Karlsruhe
Technik und Wirtschaft
UNIVERSITY OF APPLIED SCIENCES



AI Labor - Sommersemester 2020

Corona-Edition

– und –

1. Termin Computer Vision

Agenda für Heute

1. CORONA
2. Computer Vision
3. Tensorflow
4. MLflow
5. Overfitting
6. Datensätze
7. colab

CORONA

Pandemie-Modus

- › In zwei Wochen: Sprintwechsel (Remote)
 - › Review: Lösungen präsentieren
 - › Retro: Was lief gut? Was lief schlecht?
 - › Planning: Neue Aufgaben
- › In einer Woche: Remote-Betreuung
 - › Timeslots für jede Gruppe
- › Dazwischen: Slack-Workspace

Pandemie-Modus

- › Pool geschlossen → eigene Hardware / Cloud
- › eigene Hardware
 - › Docker-Setup im Git-Repository: github.com/inovex/ai-lab
 - › Macht ohne GPU keinen Spaß
- › Cloud: [Google Colab](https://colab.research.google.com/)
 - › Kostenlose GPUs und TPUs!
 - › Benötigt Google-Account (und Internet)

Gruppenfindung

1. Zu 3er Teams und ein 4er Team zusammenfinden
 - › 16 Teilnehmer → 5 Teams
2. Team-Name überlegen
3. Mitglieder und Team-Name mitteilen (slack)

Lean-Coffee

Pandemie-Modus

1. Themen sammeln
2. Stimmen auf Themen verteilen (3 Stimmen / Person)
3. Ca. 30 min Diskussion

Computer Vision

“Computer Vision is an research area that studies how to **make computers efficiently perceive, process, and understand visual data** [...]. The ultimate goal is for computers to **emulate** the striking perceptual capability of **human eyes and brains**, or even to **surpass and assist the human** in certain ways”.

– Definition von Microsoft Research

Computer Vision = Θ Computer Graphics



Bild

Computer Vision

Computer Graphics



Wissen

Computer Vision ist schwierig

MASSACHUSETTS INSTITUTE OF TECHNOLOGY

PROJECT MAC

Artificial Intelligence Group
Vision Memo. No. 100.


July 7, 1966

THE SUMMER VISION PROJECT

Seymour Papert

The summer vision project is an attempt to use our summer workers effectively in the construction of a significant part of a visual system.

Warum ist Computer Vision schwierig?




```
array([[...,  
       [0.83137256, 0.7176471, 0.6627451, 1.],  
       [0.8235294, 0.70980394, 0.654902, 1.],  
       [0.8117647, 0.7019608, 0.64705884, 1.],  
       ...  
       ],  
       [[0.8392157, 0.72156864, 0.6784314, 1.],  
       [0.8156863, 0.7058824, 0.65882355, 1.],  
       [0.8117647, 0.7019608, 0.654902, 1.],  
       ...  
       ],  
       ...  
       ], dtype=float32)
```

„Tensor“

Tensoren


- › Verallgemeinerung von Vektoren und Matrizen
 - › Tensor von Rang 0 \rightarrow Skalar
 - › Tensor von Rang 1 \rightarrow Vektor / Kovektor
 - › Tensor von Rang 2 \rightarrow Matrix
 - › Tensor von Rang 3 \rightarrow ???
- › Praktisch: mehrdimensionaler Array
 - › Rang = # Dimensionen
 - › `rank_1_tensor = np.array([1, 2, 3])`

Bilddaten als Tensoren



Graubild \rightarrow 2D Array

```
img.shape = (h, w)  
gray = img[y, x]
```

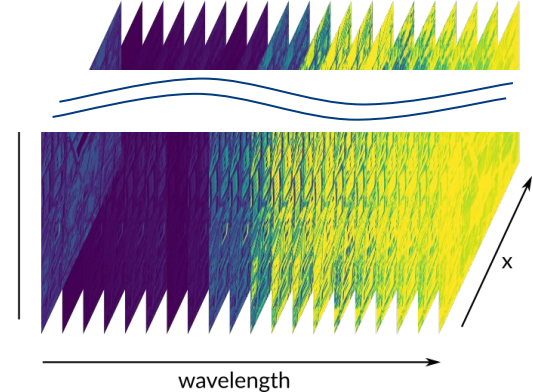


Farbbild \rightarrow 3D Array

```
img.shape = (h, w, c)  
r = img[y, x, 0]  
g = img[y, x, 1]  
b = img[y, x, 2]
```

auch:

```
img.shape = (c, h, w)
```



Spektral \rightarrow 3D Array

```
img.shape = (h, w,  $\lambda$ )  
wl_0 = img[y, x, 0]  
wl_1 = img[y, x, 1]  
wl_2 = img[y, x, 2]  
...
```

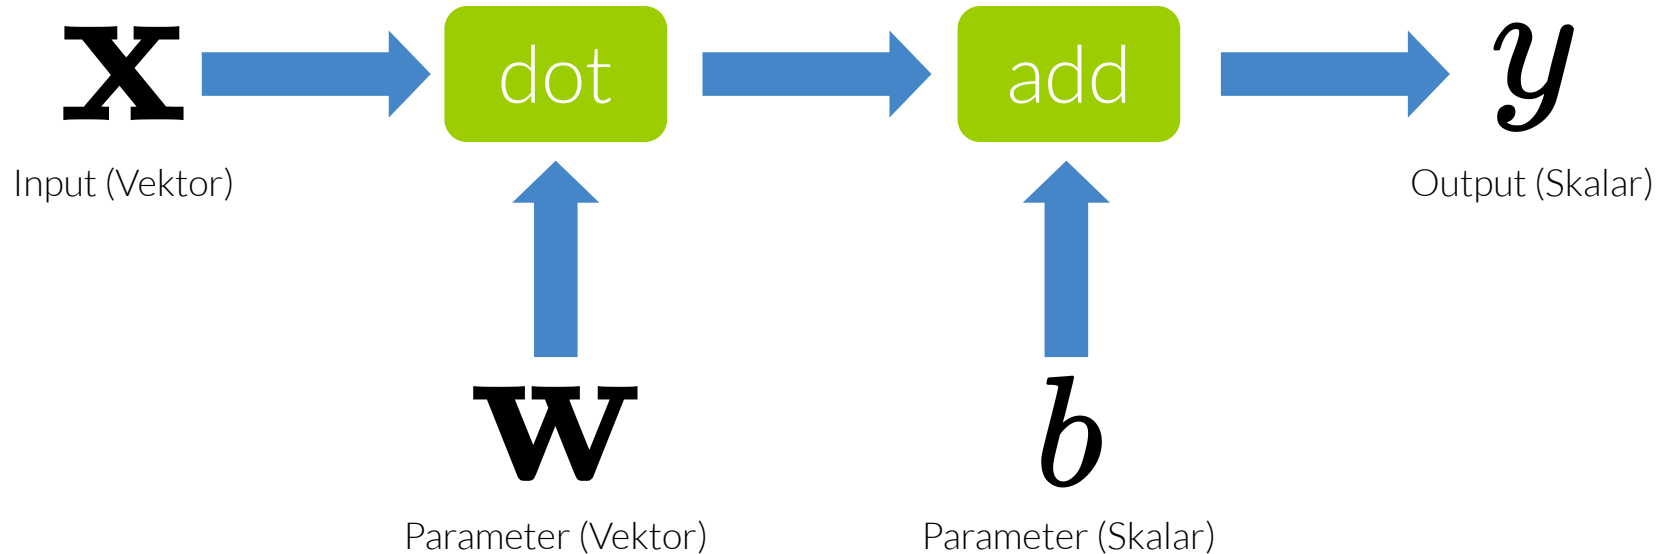


TensorFlow

- › Populäres Framework für Machine Learning
 - › Vor allem (aber nicht nur) Deep Learning
- › Kern: Graph zur Beschreibung von Datenflüssen
 - › Knoten = Transformation (add, sum, mult, ...)
 - › Kanten = Tensoren
- › Berechnung transparent auf GPU
- › Hier Fokus auf high-level Keras API

Ein einfacher Graph:

$$y = \mathbf{w}^T \mathbf{x} + b$$



```
def y(x):  
    return tf.tensordot(w, x, 1) + b
```

Hello, TensorFlow

```
from tensorflow.keras import Sequential, layers
from tensorflow.keras.losses import CategoricalCrossentropy

model = Sequential([
    layers.Dense(64, activation='relu', input_shape=32),
    layers.Dense(10)])

model.compile(optimizer=tf.keras.optimizers.Adam(0.01),
              loss=CategoricalCrossentropy(from_logits=True),
              metrics=['accuracy'])

model.fit(x=training_data, y=labels, epochs=10, batch_size=32)
predictions = model.predict(data)
```

1. Modell-Definition

```
model = Sequential([  
    layers.Dense(64, activation='relu', input_shape=32)])  
model.add(layers.Dense(10))
```

- › Sequential-API: Layer als Liste
 - › Erster Layer benötigt `input_shape`
 - › Letzter Layer definiert Modell-output
- › Mitgelieferte Layer:
 - › Core: Dense, Dropout, Flatten, Reshape, ...
 - › Convolutional: Conv1D, Conv2D, ...
 - › USW.

2. Modell-Ziel

```
model.compile(optimizer=tf.keras.optimizers.Adam(0.01),  
              loss=CategoricalCrossentropy(from_logits=True),  
              metrics=[ 'accuracy' ])
```

- › Vorbereitung für das Training
 - › Wie optimieren? → optimizer
 - › Was optimieren? → loss (und loss_weights)
 - › Wie überwachen? → metrics (und weighted_metrics)
- › Typische Losses
 - › CategoricalCrossentropy → Klassifikation
 - › MeanSquaredError → Regression

3. Modell-Training

```
model.fit(x=data, y=labels, epochs=10, batch_size=32, callbacks...  
model.fit(generator, epochs=10, batch_size=32, callbacks=[...])
```

- › Trainingsdaten: Samples + Labels oder Generator

- › Mehrerer Klassen → One-hot-encoding:

- $[1, 3, 2] \rightarrow [[1, 0, 0], [0, 0, 1], [0, 1, 0]]$

- `utils.to_categorical(class_indices, num_classes)`

- › Trainingsplan:

- › epochs, batch_size

- › callbacks (zu Beginn/Ende von Batch/Epoche)

4. Modell-Auswertung

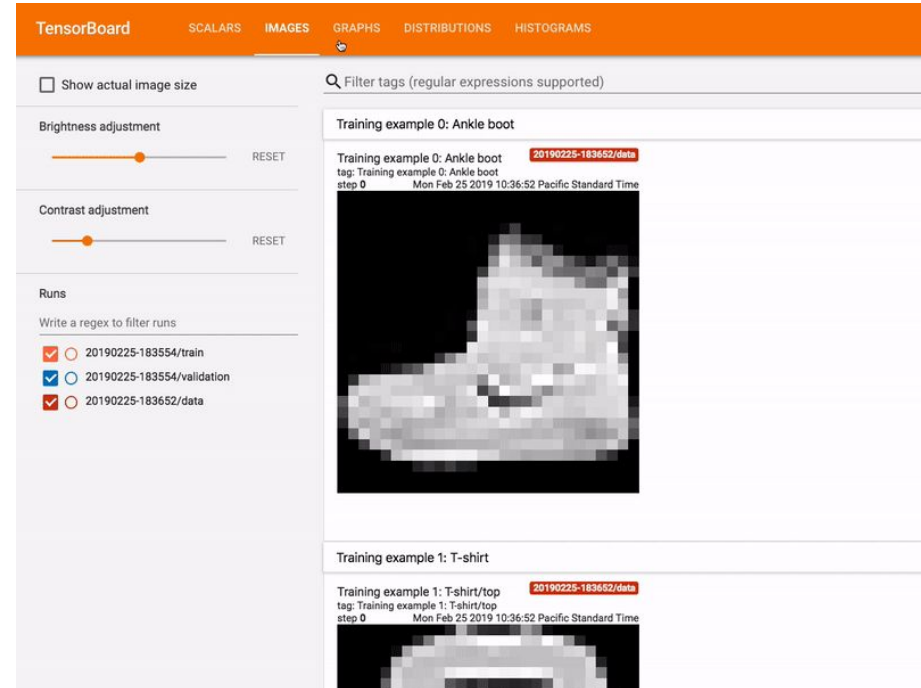
```
predictions = model.predict(data)
test_loss = model.evaluate(data, true_labels)
```

- › Wertet das Modell aus (forward-pass)
- › `predictions.shape = (data.shape[0],
 <shape-of-last-layer>)`
- › `test_loss.shape = (len(metrics),)`

TensorBoard

Visualisierung und Tooling für ML-Experimente

- › Metriken
- › Modellgraphen und -parameter
- › Trainingsdaten und (Zwischen-)Ergebnisse
- › Embeddings
- › Profiles
- › ...



mlflow™

MLFlow

Hauptfunktion

- › Experimente tracken und vergleichen
- › ML-Code verpacken
- › Modelle deployen

→ mlflow.org/docs/latest ←

Hello, MLflow

› Experimente tracken

```
from mlflow import log_metric, log_param, log_artifact

log_param('batch_size', 32)
log_param('the_answer', 42)
log_metric('accuracy', test_metrics[0])
log_artifact(path_to_the_model)
```

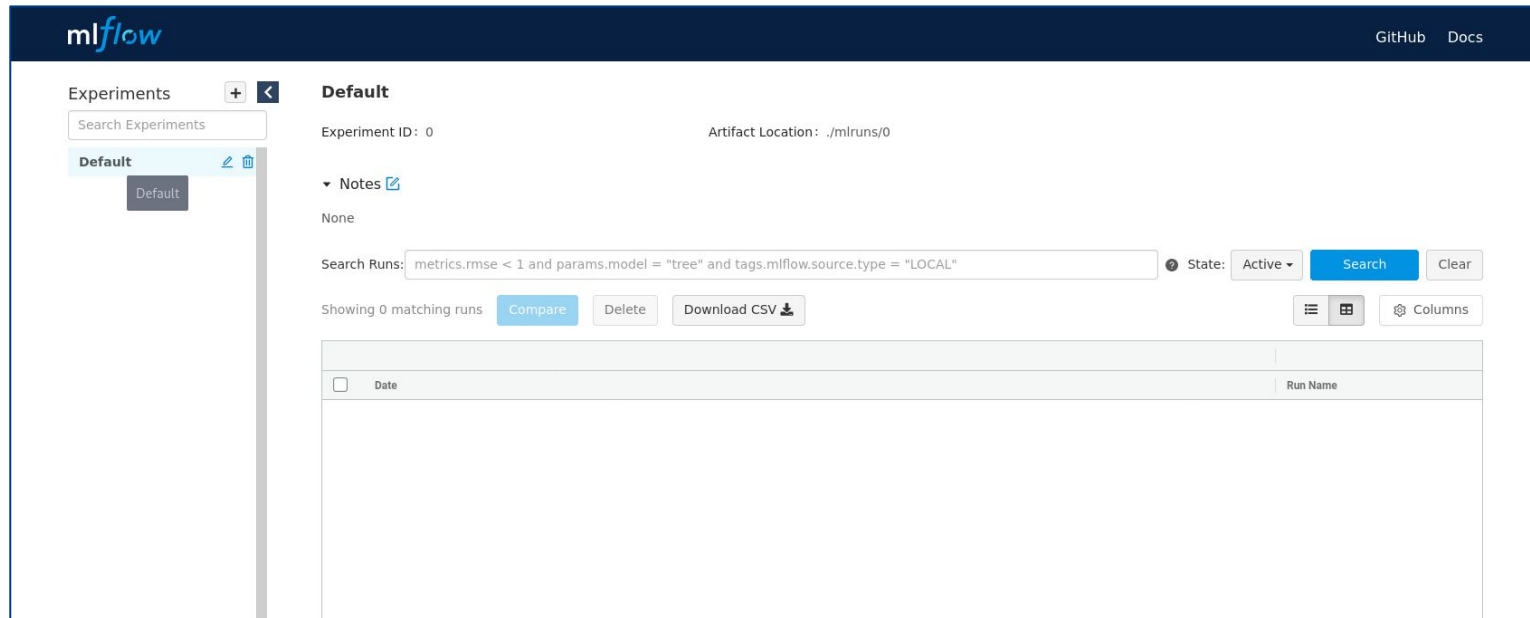
› Automatisches Logging

```
mlflow.tensorflow.autolog(every_n_iter=23)
```

Hello, MLflow


› ... und vergleichen

```
$ mlflow ui
```




Overfitting

Bias-Variance-Tradeoff




Bias-Variance-Tradeoff

Underfitting: Modell zu einfach




Bias-Variance-Tradeoff

Underfitting: Modell zu einfach




Bias-Variance-Tradeoff

Overfitting: Modell zu komplex




Bias-Variance-Tradeoff

Optimale Modellkomplexität



Bias-Variance-Tradeoff

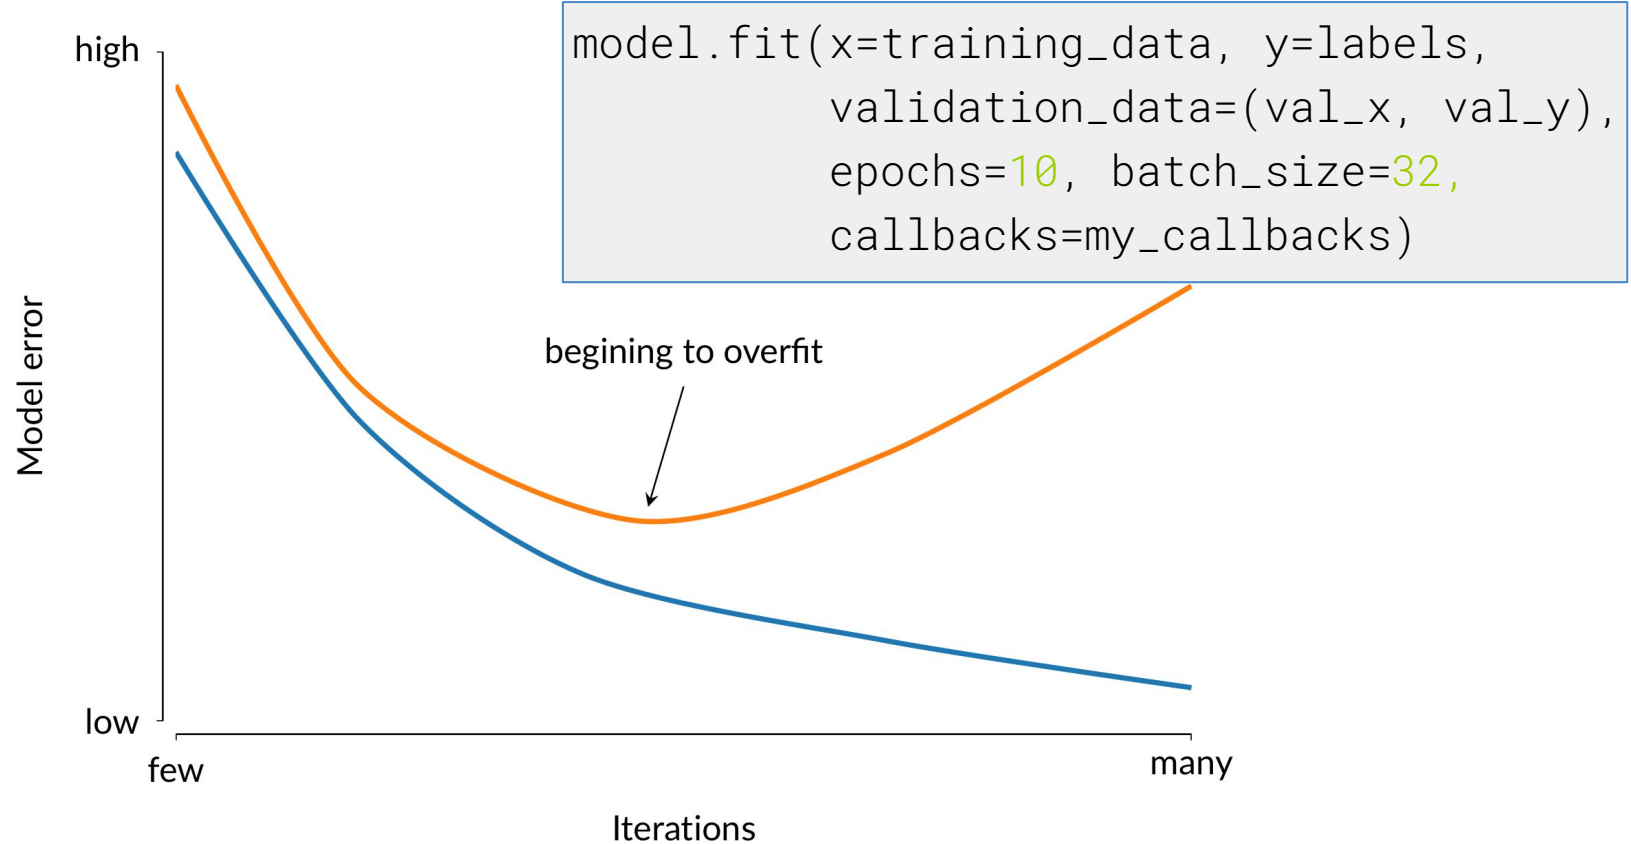
Was ist die Varianz eines Modells?



Voraussetzungen für Overfitting

- › Modell zu komplex \approx zu viele Parameter
→ Mit einfachen Modellen starten!
- › Zu lange trainiert
→ Training überwachen!
- › Zu wenig / nicht repräsentative Daten
→ Daten anschauen!

Overfitting erkennen



Datensätze


Fashion-MNIST (Zalando Research)

Abstract

We present Fashion-MNIST, a new dataset comprising of 28×28 grayscale images of 70,000 fashion products from 10 categories, with 7,000 images per category. The training set has 60,000 images and the test set has 10,000 images. Fashion-MNIST is intended to serve as a direct drop-in replacement for the original MNIST dataset for benchmarking machine learning algorithms, as it shares the same image size, data format and the structure of training and testing splits. The dataset is freely available at <https://github.com/zalando-research/fashion-mnist>.

Xiao, H., Rasul, K., & Vollgraf, R. (2017). Fashion-mnist: a novel image dataset for benchmarking machine learning algorithms.
<https://arxiv.org/pdf/1708.07747.pdf>

Fashion-MNIST (Zalando Research)



t-SNE



CIFAR-10

(Canadian Institute For Advanced Research)

“The CIFAR-10 and CIFAR-100 are **labeled subsets of the 80 million tiny images** dataset. They were collected by Alex Krizhevsky, Vinod Nair, and Geoffrey Hinton.
[...]

The CIFAR-10 dataset consists of **60000 32x32 colour images** in **10 classes**, with **6000 images per class**. There are 50000 training images and 10000 test images.”

<https://www.cs.toronto.edu/~kriz/cifar.html>

CIFAR-10

airplane




automobile



bird



cat




deer




dog




frog



horse



ship



truck



colab

Google Colab

- › In GCP gehosteter Jupyter Notebook Service
- › Kollaborative Arbeit (ähnlich GDoc) möglich
- › Frei nutzbar (inklusive GPU/TPU!)
 - › → Ressourcen sind limitiert
 - › → GPU/TPU Runtime sparsam nutzen!
- › Notebooks und Daten liegen im Google Drive



#TFDevSummit

Making the most of Colab



Feedback



<https://forms.gle/qKrigjB75GfZmw8P7>

Vielen Dank

Robin Baumann

rbaumann@inovex.de

Matthias Richter

mrichter@inovex.de

inovex GmbH

Ludwig-Erhard-Allee 6

76131 Karlsruhe

